

Multi-party DP Write-up 1

Linear Least Squares

Notations.

Let $D = [\mathbf{X}, \mathbf{y}]$ be a database of size n , with $\mathbf{X} \in \mathbb{R}^{n \times d}$ and $\mathbf{y} \in \mathbb{R}$. $[n]$ stands for the set $\{1, \dots, n\}$. We use $z \sim N(\mu, \sigma)$ to denote a variable following the Gaussian distribution with mean μ and standard deviation σ .

Linear Least Squares .

In a linear regression model the response variable is a linear function of the regressors:

$$\mathbf{y} = \mathbf{X}\theta + \epsilon \quad (1)$$

where ϵ denote the error terms.

The loss function under linear least squares model is defined as

$$L(\theta; D) = \sum_{i \in [n]} (\theta \cdot \mathbf{x}_i - \mathbf{y}_i)^2 \quad (2)$$

where \cdot is the inner product of two vectors.

Linear least squares model has a nice closed form expression for θ , such that its loss function $L(\theta; D)$ is minimized.

$$\theta = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} \quad (3)$$

Distributed Computing Problem

Let A_1, \dots, A_t and B be $t + 1$ parties (agents with computing power). Each A_i owns some of the data $\{D_j\}_{j \in S_i}$, where $S_i \subseteq [n]$. For simplicity, we assume $\{S_1, \dots, S_t\}$ partition $[n]$. The task for them is to jointly compute θ under the linear least squares model, such that the entire computation is differential-private (both intermediate output and final output should be differential private).

Privacy & Utility

There is a trade-off between privacy and utility in any mechanism.

1. We use (ϵ, δ) -differential privacy to evaluate the privacy level of our algorithm.

2. We use the concept of worst case (over input data) *expected excess risk* to evaluate the utility of our algorithm.

Definition 1. Given loss function $L(\theta; D)$ over database D , its worst case *expected excess risk* is defined as

$$E[L(\theta; D) - L(\theta^*; D)] \quad (4)$$

where θ is the output of the algorithm, θ^* is the optimal minimizer of the loss function and the randomness comes from the algorithm.

Naive Algorithm.

1. Each A_i computes $(\mathbf{C}_i, \mathbf{d}_i) = (\sum_{j \in S_i} \mathbf{x}_j \mathbf{x}_j^\top + z_1^{(d \times d)}, \sum_{j \in S_i} \mathbf{x}_j y_j + z_2^{(d)})$ and transfer it to party B , where $z_1 \sim N(0, \sigma_1)$ and $z_2 \sim N(0, \sigma_2)$.
2. Agent B computes $(\sum_{i \in [t]} \mathbf{C}_i)^{-1} (\sum_{i \in [t]} \mathbf{d}_i)$ and output it as the solution θ .

Analysis of the Naive Algorithm with $d = 1$

Definition 2. The sensitivity of function f is defined as $\Delta(f) = \sup_{(D, D')} (|f(D) - f(D')|_2)$, where (D, D') is a pair of neighboring databases.

Theorem 1 (DP with Gaussian noise). Let f be the target function, and $z \sim N(0, \sigma)$. Let $F(D) = f(D) + z$ be the output function. Fix any $\delta > 0$, we have F is (ϵ, δ) -DP, where $\epsilon = \frac{\Delta(f) \sqrt{2 \ln(2/\delta)}}{\sigma}$.

Privacy Analysis

Let $c(D) = \sum_{j \in [n]} x_j^2$ and $d(D) = \sum_{j \in [n]} x_j y_j$.

1. Each A_i 's algorithm is $(\epsilon_1 + \epsilon_2, \delta_1 + \delta_2)$ -DP, where $\epsilon_1 = \frac{\Delta(c) \sqrt{2 \ln(2/\delta_1)}}{\sigma_1}$ and $\epsilon_2 = \frac{\Delta(d) \sqrt{2 \ln(2/\delta_2)}}{\sigma_2}$.

When $d = 1$, each A_i computes and outputs $(c_i, d_i) = (\sum_{j \in S_i} x_j^2 + z_1, \sum_{j \in S_i} x_j y_j + z_2)$, where $z_1 \sim N(0, \sigma_1)$ and $z_2 \sim N(0, \sigma_2)$. It is obvious that $\Delta(c_i) = \Delta(c)$ and $\Delta(d_i) = \Delta(d)$ for all $i \in [t]$.

Fix $\delta_1 > 0$, c_i is (ϵ_1, δ_1) -DP, where $\epsilon_1 = \frac{\Delta(c) \sqrt{2 \ln(2/\delta_1)}}{\sigma_1}$. Similarly, d_i is (ϵ_2, δ_2) -DP, where $\epsilon_2 = \frac{\Delta(d) \sqrt{2 \ln(2/\delta_2)}}{\sigma_2}$. By the composition theorem of differential privacy, each A_i 's algorithm is $(\epsilon_1 + \epsilon_2, \delta_1 + \delta_2)$ -DP.

2. B 's algorithm.

It is easy to see B 's output is

$$\theta = \frac{\sum_{i \in [n]} x_i y_i + t z_2}{\sum_{i \in [n]} x_i^2 + t z_1}. \quad (5)$$

$$= \frac{d(D)}{c(D) + t z_1} + \frac{t z_2}{c(D) + t z_1}. \quad (6)$$

We only consider the case when $(X^\top X)$ is invertible. Here, we assume WLOG that $c(D) = \sum_{i \in [n]} x_i^2 \geq b_l > 0$. We treat $\frac{d(D)}{c(D) + t z_1}$ as our target function and $\frac{t z_2}{c(D) + t z_1}$ as the noise.

Pick b_z such that $0 < b_z < b_l$, let $1 - \delta_z$ be the probability that $t z_1 \in [-b_z, b_z]$. Then, with probability $1 - \delta_z$, $c(D) + t z_1 \geq b_l - b_z$, and therefore $\Delta(\frac{d(D)}{c(D) + t z_1}) \leq \Delta(d)/(b_l - b_z)$.

We need to separate the discussion on the upper bound of $c(D)$.

- (a) If $c(D)$ does not have an upper bound, $\frac{t z_2}{c(D) + t z_1}$ could be 0 and B 's algorithm does not have privacy guarantee. We need to modify B 's algorithm by adding $z_3 \sim N(0, \sigma_3)$ to θ .
- (b) If $c(D) \leq b_u$. Then, with probability $1 - \delta_z$, $c(D) + t z_1 \leq b_u + b_z$ and therefore $\frac{t z_2}{c(D) + t z_1} \geq \frac{t z_2}{b_u + b_z}$. It is obvious that $\frac{t z_2}{b_u + b_z} \sim N(0, \frac{t}{b_u + b_z} \sigma_2)$.

Fix $\delta_3 > 0$, we have B 's algorithm is $(\epsilon_3, \delta_z + \delta_3)$ -DP, where $\epsilon_3 \leq \frac{(\Delta(d)/(b_l - b_z))\sqrt{2 \ln(2/\delta_3)}}{\frac{t}{b_u + b_z} \sigma_2} = \frac{\Delta(d)(b_u + b_z)\sqrt{2 \ln(2/\delta_3)}}{(b_l - b_z)t \sigma_2}$

Utility Analysis

We calculate $E[L(\theta; D) - L(\theta^*; D)]$, where θ^* is the optimal minimizer to the loss function. In the following discussion, we assume $c(D) = \sum_{i \in [n]} x_i^2$ has some upper bound ($c(D) \leq b_u$), and therefore B 's algorithm has certain level of privacy guarantee. In the case when $c(D)$ does not have an upper bound, it is easy to see that adding another $z_3 \sim N(0, \sigma_3)$ to θ will bring extra $\sigma_3^2 c(D)$ to the expected excess risk.

It is known that $\theta^* = \frac{d(D)}{c(D)}$ under the linear least model.

$$E[L(\theta; D) - L(\theta^*; D)] = E[\sum_{i \in [n]} (\frac{d(D) + t z_2}{c(D) + t z_1} x_i - y_i)^2 - (\frac{d(D)}{c(D)} x_i - y_i)^2] \quad (7)$$

We focus on its i -th term.

$$E[(\frac{d(D) + tz_2}{c(D) + tz_1}x_i - y_i)^2 - (\frac{d(D)}{c(D)}x_i - y_i)^2] \quad (8)$$

$$= E[(\frac{(d(D) + tz_2)^2}{(c(D) + tz_1)^2}x_i^2 + y_i^2 - 2\frac{d(D) + tz_2}{c(D) + tz_1}x_i y_i) - (\frac{d(D)^2}{c(D)^2}x_i^2 + y_i^2 - 2\frac{d(D)}{c(D)}x_i y_i)] \quad (9)$$

$$= E[(\frac{(d(D) + tz_2)^2}{(c(D) + tz_1)^2}x_i^2] - \frac{d(D)^2}{c(D)^2}x_i^2 \quad (10)$$

$$= E[(\frac{(d(D)^2 + t^2 z_2^2 + 2d(D)tz_2)}{(c(D)^2 + t^2 z_1^2 + 2c(D)tz_1)}x_i^2] - \frac{d(D)^2}{c(D)^2}x_i^2 \quad (11)$$

$$= (\frac{d(D)^2 + t^2 \sigma_2^2}{c(D)^2 + t^2 \sigma_1^2} - \frac{d(D)^2}{c(D)^2})x_i^2 \quad (12)$$

Hence, we have $E[L(\theta; D) - L(\theta^*; D)] = (\frac{d(D)^2 + t^2 \sigma_2^2}{c(D)^2 + t^2 \sigma_1^2} - \frac{d(D)^2}{c(D)^2})c(D)$.

Algorithm based on Stochastic gradient descent

1. B chooses θ_0 from the convex set \mathcal{C} , learning rate η and a permutation π over n copies of $[n]$. B broadcasts (θ_0, η, π)
2. In each iteration i , the agent A_j that owns data $D_{\pi(i)}$ will compute $\theta_{i+1} = \Pi_{\mathcal{C}}[\theta_i - \eta(i)(n\Delta(l(\theta_i; D_{\pi(i)})) + z_i)]$, where z_i is a zero mean Gaussian noise and $\Pi_{\mathcal{C}}$ is a projection function to set \mathcal{C} . A_j broadcasts $G_i = n\Delta(l(\theta_i; D_{\pi(i)}))$.
3. Each of the remaining A_k will update $\theta_{i+1} = \Pi_{\mathcal{C}}[\theta_i - \eta(i)G_i]$